

# A Introduction to The Center for Advanced Computing

Brock Palen  
brockp@umich.edu

TBD

# Outline

- 1 Resources
  - Hardware
  - Default Software
- 2 Mechanics: Usage
  - Compiling programs
  - The Batch System
- 3 The Scheduler
  - Understanding the Scheduler
  - Scheduler Commands
- 4 Summary
  - Resources and Access
  - Job Management
  - Contact

# Hardware

## Compute Hardware

- 1 Altix node, 32 cores
- 468 Opteron nodes, over 1224 cores
- 400+ nodes on CAEN Grid [update this](#)
- Gigabit networking, Myrinet networking, Infiniband networking
- Upto 96GB of memory (64GB public) for SMP work

# Hardware: nyx

## Nyx

- nyx is the Opteron cluster;
- Login: `nyx-login.engin.umich.edu`
- Currently has 6TB NFS file system
- Running RedHat Enterprise Linux 4

# Hardware: bighouse

## Bighouse

- bighouse is our Itanium SMP machine;
- Login: `bighouse.engin.umich.edu`
- Shares nyx's 6TB NFS file system
- Running SUSE Linux Enterprise Server 10
- ProPack 5 from SGI
- Only available for benchmarking (Private)

# Hardware: Grid

## CAEN Grid

- 400+ Nodes, Some Dual Core
- All nodes have 2GB Ram
- FAST Single CPU's
- Some Parallel Ability
- Short Jobs Only

# Software

## Nyx Defaults

- OpenMPI
- PGI Compilers
- PBS commands

## Bighouse Defaults

- Message Passing Toolkit (MPT)
- Intel Compilers
- PBS commands

# Manipulating Software

All CAC systems use modules to control software. Users *can* and *should* write their own modules if needed.

## module commands

- `module list`  
Show loaded modules
- `module load modulename`  
Load *modulename* for use
- `module avail modulename`  
Show available versions of module *modulename*
- `module rm modulename`  
Remove currently loaded module



# Module Fun

## Module Customization

- `/privatemodules/default`  
Allows users to change their default modules.
- `/privatemodules/module/version`  
Holds user created module
- `man modulefile`

# Module Example

## Module Example

- List Loaded Modules  
`module list`
- Show All Modules  
`module avail`
- Show All Versions of openmpi  
`module avail openmpi`
- Load FFTW  
`module load fftw`
- Show Variables defined by FFTW  
`module show fftw`  
`echo $FFTW_LINK`

# Module Example

## Module Example

- List Loaded Modules  
`module list`
- Show All Modules  
`module avail`
- Show All Versions of openmpi  
`module avail openmpi`
- Load FFTW  
`module load fftw`
- Show Variables defined by FFTW  
`module show fftw`  
`echo $FFTW_LINK`

# Module Example

## Module Example

- List Loaded Modules  
`module list`
- Show All Modules  
`module avail`
- Show All Versions of openmpi  
`module avail openmpi`
- Load FFTW  
`module load fftw`
- Show Variables defined by FFTW  
`module show fftw`  
`echo $FFTW_LINK`

# Module Example

## Module Example

- List Loaded Modules  
`module list`
- Show All Modules  
`module avail`
- Show All Versions of openmpi  
`module avail openmpi`
- Load FFTW  
`module load fftw`
- Show Variables defined by FFTW  
`module show fftw`  
`echo $FFTW_LINK`

# Module Example

## Module Example

- List Loaded Modules  
`module list`
- Show All Modules  
`module avail`
- Show All Versions of openmpi  
`module avail openmpi`
- Load FFTW  
`module load fftw`
- Show Variables defined by FFTW  
`module show fftw`  
`echo $FFTW_LINK`

# Tools

- All of the standard GNU/Linux tools are also available: make, autoconf, awk, sed, Perl, Python,
- We support emacs, vi{m}, and nano (a pico-like editor) on the clusters. etc.
- Only use notepad on Windows!
- If made on windows fix with `dos2unix filename`

# Tools

- All of the standard GNU/Linux tools are also available: make, autoconf, awk, sed, Perl, Python,
- We support emacs, vi{m}, and nano (a pico-like editor) on the clusters. etc.
- Only use notepad on Windows!
- If made on windows fix with `dos2unix filename`



# Tools

- All of the standard GNU/Linux tools are also available: make, autoconf, awk, sed, Perl, Python,
- We support emacs, vi{m}, and nano (a pico-like editor) on the clusters. etc.
- Only use notepad on Windows!
- If made on windows fix with `dos2unix filename`

# Compile Code

## Nyx

- Use: `mpicc`, `mpiCC`, `mpif90` for MPI code
- Use: `pgcc`, `pgCC`, `pgf90` with `-mp` for OpenMP Code

## Bighouse

- Use: `icc`, `icpc`, `ifort` with `-lmpi` for MPI code
- Use: `icc`, `icpc`, `ifort` with `-openmp` for OpenMP code

## CAEN Grid

TBD

# Compile Example

```
Copy Code: cp ~brockp/mpicodes.tar ~
tar -xvf mpicodes.tar
cd ~/mpicodes
```

## Serial Code

- Fortran 90

```
pgf90 -fastsse -o f90hello helloworld.f90
```

- C

```
pgcc -fastsse -o chello helloworld.c
```

# Compile Example Ctd.

## MPI Code

- `make`
- `mpirun -np 2 c_ex01`
- Thats it... Ok not really
- `make clean`
- `mpicc -fastsse -o c_ex01 c_ex01.c`

# Compile Example Ctd.

## MPI Code

- make
- mpirun -np 2 c\_ex01
- Thats it... Ok not really
- make clean
- mpicc -fastsse -o c\_ex01 c\_ex01.c

# Compile Example Ctd.

## MPI Code

- `make`
- `mpirun -np 2 c_ex01`
- Thats it... Ok not really
- `make clean`
- `mpicc -fastsse -o c_ex01 c_ex01.c`

# Introduction to the PBS Batch System

## PBS

- All access to the compute nodes (everything other than the login node) is via the batch system
- We use a system called Torque, it is derived from PBS
- The batch system controls access to queues
- The scheduling system decides if and where jobs can run
- There is a single public queue: cac
- There are many private queues for people who own or rent nodes
- If you don't know use the route queue

# Introduction to the PBS Batch System

## PBS

- All access to the compute nodes (everything other than the login node) is via the batch system
- We use a system called Torque, it is derived from PBS
- The batch system controls access to queues
- The scheduling system decides if and where jobs can run
- There is a single public queue: cac
- There are many private queues for people who own or rent nodes
- If you don't know use the route queue



# Introduction to the PBS Batch System

## PBS Files

The steps to using the batch system are:

- 1 Create a batch file: this is a short (5-15 lines) text file with some batch commands and the commands to run your program
- 2 Submit the file to the batch system
- 3 Check on the status of your job
- 4 Delete your job if you want to cancel it

# Creating a PBS Batch File

A simple single cpu example

```
#!/bin/sh
#PBS -N 1-cpu
#PBS -l nodes=1,walltime=1:00:00
#PBS -m abe
#PBS -M brockp@umich.edu
#PBS -q route
#PBS -j oe
#PBS -V
cat $PBS_NODEFILE
cd ~/input1dir/
mcnp5.mpi i=input o=output r=restart
```

# Creating a PBS Batch File

A simple single cpu example

```
#!/bin/sh
#PBS -N 1-cpu
#PBS -l nodes=1,walltime=1:00:00
#PBS -m abe
#PBS -M brockp@umich.edu
#PBS -q route
#PBS -j oe
#PBS -V
cat $PBS_NODEFILE
cd ~/input1dir/
mcnp5.mpi i=input o=output r=restart
```

# Creating a PBS Batch File

A simple single cpu example

```
#!/bin/sh
#PBS -N 1-cpu
#PBS -l nodes=1,walltime=1:00:00
#PBS -m abe
#PBS -M brockp@umich.edu
#PBS -q route
#PBS -j oe
#PBS -V
cat $PBS_NODEFILE
cd ~/input1dir/
mcnp5.mpi i=input o=output r=restart
```

# Creating a PBS Batch File

A simple single cpu example

```
#!/bin/sh
#PBS -N 1-cpu
#PBS -l nodes=1,walltime=1:00:00
#PBS -m abe
#PBS -M brockp@umich.edu
#PBS -q route
#PBS -j oe
#PBS -V
cat $PBS_NODEFILE
cd ~/input1dir/
mcnp5.mpi i=input o=output r=restart
```

# Creating a PBS Batch File

A simple single cpu example

```
#!/bin/sh
#PBS -N 1-cpu
#PBS -l nodes=1,walltime=1:00:00
#PBS -m abe
#PBS -M brockp@umich.edu
#PBS -q route
#PBS -j oe
#PBS -V
cat $PBS_NODEFILE
cd ~/input1dir/
mcnp5.mpi i=input o=output r=restart
```

# Creating a PBS Batch File

A simple single cpu example

```
#!/bin/sh
#PBS -N 1-cpu
#PBS -l nodes=1,walltime=1:00:00
#PBS -m abe
#PBS -M brockp@umich.edu
#PBS -q route
#PBS -j oe
#PBS -V

cat $PBS_NODEFILE
cd ~/input1dir/
mcnp5.mpi i=input o=output r=restart
```

# Creating a PBS Batch File

A simple single cpu example

```
#!/bin/sh
#PBS -N 1-cpu
#PBS -l nodes=1,walltime=1:00:00
#PBS -m abe
#PBS -M brockp@umich.edu
#PBS -q route
#PBS -j oe
#PBS -V
cat $PBS_NODEFILE
cd ~/input1dir/
mcnp5.mpi i=input o=output r=restart
```



# Creating a PBS Batch File

A simple single cpu example

```
#!/bin/sh
#PBS -N 1-cpu
#PBS -l nodes=1,walltime=1:00:00
#PBS -m abe
#PBS -M brockp@umich.edu
#PBS -q route
#PBS -j oe
#PBS -V
cat $PBS_NODEFILE
cd ~/input1dir/
mcnp5.mpi i=input o=output r=restart
```

# Creating a PBS Batch File

A more complicated example:

```
#!/bin/sh
#PBS -N mcnp-8x2
#PBS -l nodes=8:ppn=2,walltime=8:00:00
#PBS -q route
#PBS -M brockp@umich.edu
#PBS -m ae
#PBS -j oe
#PBS -V
cd ${HOME}/input2/
echo "I ran on: "
cat $PBS_NODEFILE
mpirun -np 16 mcnp5.mpi i=input2 o=output2 r=restart2
```

# Submitting, Checking, and Deleting Batch Jobs

- After you create your PBS script, you need to submit it:

```
$ qsub mcnp.q
542.nyx-login.engin.umich.edu
```

- After you submit your script, you can check on the status of your job:

```
$ qstat -au brockp
nyx-login.engin.umich.edu:
Job ID          Username Queue   Jobname   SessID NDS   TSK Memory Time   S Time
-----
542.nyx-login.engin. brockp  short    mcnp-8x2   18922    8  --    --   08:00 R 00:00

$ checkjob 542
[... lots of output ...]
```

- If you want to delete your job:

```
$ qdel 542
```

# Submitting, Checking, and Deleting Batch Jobs

- After you create your PBS script, you need to submit it:

```
$ qsub mcnp.q
542.nyx-login.engin.umich.edu
```

- After you submit your script, you can check on the status of your job:

```
$ qstat -au brockp
nyx-login.engin.umich.edu:
Job ID          Username Queue   Jobname   SessID NDS   TSK Memory Time   S Time
-----
542.nyx-login.engin. brockp  short    mcnp-8x2   18922    8  --    --   08:00 R 00:00

$ checkjob 542
[... lots of output ...]
```

- If you want to delete your job:

```
$ qdel 542
```

# Submitting, Checking, and Deleting Batch Jobs

- After you create your PBS script, you need to submit it:

```
$ qsub mcnp.q
542.nyx-login.engin.umich.edu
```

- After you submit your script, you can check on the status of your job:

```
$ qstat -au brockp
nyx-login.engin.umich.edu:
Job ID          Username Queue   Jobname   SessID NDS   TSK Memory Time   S Time
-----
542.nyx-login.engin. brockp  short   mcnp-8x2   18922    8  --    --   08:00 R 00:00

$ checkjob 542
[... lots of output ...]
```

- If you want to delete your job:

```
$ qdel 542
```

# PBS Email

PBS will send an email at the start and end of your job if you use the `-m` and `-M` options in your PBS script. The email after a job completes successfully looks like:

```
Date: Sun, 30 Apr 2006 12:50:17 -0400
From: adm <adm@nyx-login.engin.umich.edu>
To: "Palen, Brock E" <brockp@umich.edu>
Subject: PBS JOB 542.nyx-login.engin.umich.edu
-----
```

```
PBS Job Id: 542.nyx-login.engin.umich.edu
Job Name:   mcnp-8x2
Execution terminated
Exit_status=0
resources_used.cput=13:17:26
resources_used.mem=1220672kb
resources_used.vmem=11146704kb
resources_used.walltime=00:49:57
```

# PBS Example

## PBS Example Script

```
cd /mpicodes
nano run
Edit #PBS -M
Ctl+o
Ctl+x
qsub run
```

# Understanding the Scheduler

The scheduler determines what jobs can run, when the can run, and where. There are many factors that go into the scheduler's decision.

- Limits

- Maximum number jobs eligible for scheduling: 4
- Maximum number of CPUs in use by one person: depends on queue
- Maximum number of jobs in the queue at one time: no limit

- Priority

- Who you are: user and group level priorities
- How long you've waited: the longer you wait, the higher your priority
- Your recent usage (fairshare): People with less usage over the past month will have a higher priority than those with a lot of usage



# Understanding the Scheduler

The scheduler determines what jobs can run, when the can run, and where. There are many factors that go into the scheduler's decision.

- Limits
  - Maximum number jobs eligible for scheduling: 4
  - Maximum number of CPUs in use by one person: depends on queue
  - Maximum number of jobs in the queue at one time: no limit
- Priority
  - Who you are: user and group level priorities
  - How long you've waited: the longer you wait, the higher your priority
  - Your recent usage (fairshare): People with less usage over the past month will have a higher priority than those with a lot of usage


# Understanding the Scheduler

- Reservations

- Advance reservations: holds nodes for users or groups
- Job reservations: scheduler will reserve nodes for the next several jobs in each queue


- Backfill

- If the reservations leave holes in the schedule, they may be filled by short jobs that otherwise would have waited.



# Understanding the Scheduler

- Reservations
  - Advance reservations: holds nodes for users or groups
  - Job reservations: scheduler will reserve nodes for the next several jobs in each queue
- Backfill
  - If the reservations leave holes in the schedule, they may be filled by short jobs that otherwise would have waited.



# Understanding the Scheduler

There are several commands that can give you insight into the scheduler's decisions.

- `showq` — shows the state of the queue at that moment in time, showing the running jobs in order of soonest to finish to longest to finish; the idle jobs in order of priority; and the blocked jobs in the order they were submitted
- `diagnose -p` — shows the factors that go into computing the priority for all of the idle jobs
- `checkjob jobnumber` — for idle jobs this will show why the job can't start
- `showstart jobnumber` — this makes a (poor) estimate of when the job will start

# Summary

- Resources
  - Lots of CPUs
  - A reasonable amount of software
  - Watch or subscribe to <http://cac.engin.umich.edu> for updates
- Access
  - All access is via the SSH family of commands: `ssh`, `sftp`, `scp`
  - There are lots of clients for these commands for the different platforms
  - There is no graphical access, everything is via the command line

# Summary

- Job Submission
  - Every job needs a PBS script file
  - Two most important commands: `qsub` and `qstat -au username`
- Job Scheduling
  - Scheduling depends on a lot of factors, it is best to submit jobs and let the scheduler optimize for their start.

# Summary

- News: <http://cac.engin.umich.edu>
  - RSS feed
  - New of changes, outages, other pertinent piece of information
- Contact: [cac-support@umich.edu](mailto:cac-support@umich.edu)
  - Questions or concerns should be sent here (not to an individual) since this is read by six people. The odds of a quick reply are best this way.
  - We aren't parallel programmers, but we'll do what we can to help.

# Example

Change example to

example

Open a shell....

- 1 `cp -r ~brockp/mcnp_example /`
- 2 `cat mcnp.q`
- 3 `module load mcnp5`
- 4 `qsub mcnp.q`
- 5 `qstat -u $USER`



# Example

Change example to

example

Open a shell....

- 1 `cp -r ~brockp/mcnp_example /`
- 2 `cat mcnp.q`
- 3 `module load mcnp5`
- 4 `qsub mcnp.q`
- 5 `qstat -u $USER`

# Example

Change example to

example

Open a shell....

- 1 `cp -r ~brockp/mcnp_example /`
- 2 `cat mcnp.q`
- 3 `module load mcnp5`
- 4 `qsub mcnp.q`
- 5 `qstat -u $USER`

# Example

Change example to

example

Open a shell....

- ① `cp -r ~brockp/mcnp_example /`
- ② `cat mcnp.q`
- ③ `module load mcnp5`
- ④ `qsub mcnp.q`
- ⑤ `qstat -u $USER`

# Example

Change example to

example

Open a shell....

- ① `cp -r ~brockp/mcnp_example /`
- ② `cat mcnp.q`
- ③ `module load mcnp5`
- ④ `qsub mcnp.q`
- ⑤ `qstat -u $USER`